

✉ khaykingleb@gmail.com

⌚ https://github.com/khaykingleb

.linkedin.com/in/khaykingleb

GLEB KHAYKIN

📍 Amsterdam, The Netherlands

+31 06 26172415 ☎

https://khaykingleb.com 🌐

https://t.me/blog_khaykingleb 📲

SUMMARY

ML Infrastructure / Platform Engineer with 4+ years of experience architecting high-performance ML systems, multi-cluster platforms, and large-scale inference/training pipelines. Specialized in distributed systems, Kubernetes, Terraform, observability, and ML platform automation. Delivered 50%+ GPU cost reductions, 4x CI/CD improvements, and inference systems serving 60k+ req/hour.

EXPERIENCE

Together AI

Senior ML Platform Engineer

Amsterdam, Netherlands

July 2025 — Present

- Migrated model development workloads off a saturated shared EKS to a dedicated one with Karpenter, eliminating 1h manual node scaling
- Engineered model sync pipeline and integration tests for 90+ models, unlocking dedicated and serverless inference for fine-tuned models
- Unified observability across 6+ K8s clusters into a single Grafana Cloud stack and shipped 30+ IaC-defined dashboards
- Accelerated CI/CD pipelines 4x through BuildKit mount caching, Docker optimization, and self-hosted runners with shared cache

Stealth Startup

MLOps/DevOps Engineer

Remote

July 2023 — July 2025

- Provisioned and operated hybrid-cloud K3s GPU clusters with Tailscale-meshed networking, reducing GPU costs by >50%
- Standardized K8s monitoring and alerting across 4 clusters in Grafana Cloud, reducing issue detection time from ad hoc to under 1 hour
- Optimized ASR/diarization with NVIDIA Triton BLS and model selection, achieving 11x speedup and cutting processing costs by 5x
- Delivered a multi-stage financial document parser using vision LLMs, reaching 70% automation and securing a venture fund as first client

Myna Labs

ML/MLOps Engineer

Remote

November 2022 — March 2024

- Owned a priority-queued speech inference API (TTS/VC/ASR) on FastAPI and NVIDIA Triton, serving 60k+ req/hour
- Built a voice cloning service using LoRA fine-tuning, reducing voice addition time from days to under 10 minutes
- Modernized and automated speech-related infrastructure by migrating from Docker Compose to Terraform-managed GKE with autoscaling

Stealth Startup

Data Scientist

Remote

September 2021 — July 2022

- Designed distributed PySpark pipeline for financial sentiment analysis, scraping and processing S&P 500 news data with BERT models
- Developed learning-to-rank system for quantitative stock selection in long-short trading strategies
- Engineered automated ETL workflows on Databricks to aggregate and analyze data from 150+ financial API endpoints

PROJECTS

- **Research Playground:** Production-ready PyTorch Lightning framework with Hydra configs and distributed training on self-hosted K3s
- **Dotfiles:** Declarative system configuration with Nix flakes and home-manager for reproducible macOS environments
- **Huffman Archiver:** C++ implementation of Huffman coding for lossless data compression with Conan package management
- **Personal Website:** Full-stack blog and portfolio with Next.js frontend, Supabase backend, and live content sync via Notion API

EDUCATION

NRU “Higher School of Economics”

Moscow, Russia

B.S., Computer Science and Finance (summa cum laude)

September 2018 — July 2022

CFA Institute

Moscow, Russia

Level 1 passed

February 2021

SKILLS

- **Languages:** Go, Python, Rust, C++, SQL, TypeScript
- **DevOps:** Kubernetes, Terraform, AWS, GCP, Docker, Helm, ArgoCD, Ansible, Nix, Grafana, Prometheus, Loki
- **MLOps:** NVIDIA Triton, TensorRT, ONNX, Volcano, Flyte, DVC
- **R&D:** PyTorch, Lightning, Hydra, W&B, NumPy, Pandas, SciPy
- **Data Engineering:** Spark, Polars, Kafka, RabbitMQ, Celery
- **Storage:** PostgreSQL, Redis, MongoDB, S3/R2, PGVector
- **Frontend:** React, Next.js, Tailwind