khaykingleb@gmail.com
https://github.com/khaykingleb
https://linkedin.com/in/khaykingleb

Gleb Khaykin

Amsterdam, The Netherlands

+31 06 26172415
https://khaykingleb.com
https://t.me/blog_khaykingleb

Summary

Senior Infrastructure/Platform Engineer specializing in distributed systems, GPU cluster infrastructure, and Kubernetes platforms, with 5+ years across AI startups. Designed and own Together AI's RL fine-tuning platform end-to-end, from distributed training orchestration and multi-cluster GPU scheduling to autoscaling and Weka-backed storage on self-managed GPU clusters. Cut GPU costs by 50%+, unified observability across 9+ K8s clusters, and shipped 60k+ req/hour production inference.

Experience

Together AI — Amsterdam, Netherlands
Senior Platform Engineer, Research Platform — *July 2025 — Present*

- Designed and now own the RL fine-tuning platform end-to-end, from gRPC/HTTP API and Kafka-driven orchestration to Volcano-based GPU scheduling and Weka-backed storage, supporting multi-week distributed training runs across up to 256 GPUs (32 nodes)
- Adopted Meta's Monarch actor framework into the RL platform, working closely with the PyTorch team on the Kubernetes integration, now powering internal multi-node GPU training and inference workflows for 7 research engineers
- Architected the platform's multi-tier capacity model with best-fit placement across dedicated, reserved, and shared GPU pools, with the latter two autoscaled to demand, eliminating manual capacity planning per customer tier
- Migrated platform workloads from a click-ops shared cluster to a Terraform-based, Karpenter-autoscaled EKS cluster, eliminating 1-hour manual node scaling and supporting 30+ concurrent per-PR preview environments via ArgoCD ApplicationSet generators
- Engineered model sync and integration tests for 90+ fine-tunable models, unlocking dedicated and serverless inference for fine-tuned variants
- Consolidated the team's observability from a fragmented setup into a unified Grafana Cloud stack with IaC-managed dashboards and alerting across 9+ K8s clusters, with new clusters now onboardable in minutes
- Modernized CI/CD, achieving a 4x speedup with BuildKit optimizations and delivering self-hosted GPU-based CI for 90+ engineers via ARC

Stealth AI Startup — Remote
MLOps/DevOps Engineer, AI Platform — *July 2023 — July 2025*

- Built a Terraform-managed platform across 4 environments, combining AWS EKS with Tailscale Operator-connected K3s GPU clusters and ArgoCD-deployed workloads, reducing GPU costs by >50% compared to hyperscaler pricing
- Established declarative K8s monitoring and alerting in Grafana Cloud across 3 clusters, cutting issue detection from ad hoc to under 1 hour
- Optimized a speech inference pipeline using WhisperX for ASR, NeMo for diarization, and Demucs for denoising on NVIDIA Triton with per-request model selection, achieving 11x speedup and 5x cost reduction on 1-hour audio inputs
- Delivered a multi-stage financial document parser using vision LLMs and LangGraph agents on an async Celery pipeline, achieving 70% end-to-end automation and landing a venture fund as the first client

Myna Labs (from the creators of Snap Lenses and Cameos) — Remote
MLOps Engineer, Speech Synthesis — *November 2022 — March 2024*

- Architected and owned a speech inference platform on FastAPI, Celery, and NVIDIA Triton with per-domain priority queues for text-to-speech, voice conversion, and ASR, serving mobile and web apps at 60k+ req/hour
- Shipped a voice cloning service on Tortoise TTS using LoRA fine-tuning with precomputed speaker latents, cutting per-voice onboarding time from days to under 10 minutes, enabling custom voices in production
- Replatformed from Docker Compose to Terraform-provisioned GKE with an elastic GPU node pool across staging and production

Stealth Finance Startup — Remote
Data Scientist, Quant Finance — *September 2021 — July 2022*

- Built a distributed PySpark pipeline that transformed S&P 500 news and analyst reports into daily company-level FinBERT sentiment features
- Developed ListNet-based learning-to-rank models for long-short equity selection, evaluated with nDCG, Sharpe, and Sortino

Education

NRU "Higher School of Economics" — Moscow, Russia
B.S., Computer Science and Finance, summa cum laude — *September 2018 — July 2022*

- Ranked in the top 3% of the cohort
- Passed CFA Level 1 in February 2021, scoring in the top 10% of candidates worldwide

Skills

- **Languages**: Go, Python, Rust, C++, SQL, TypeScript
- **Infrastructure**: Kubernetes, EKS/GKE/K3s, Karpenter, Terraform, Ansible, MAAS, NetBox, AWS/GCP, Docker, Helm, ArgoCD, ARC, Nix
- **Networking & Security**: Tailscale, Cloudflare, InfiniBand/UFM, Dex/OIDC, TLS/mTLS, IAM/RBAC, NetworkPolicy, External Secrets
- **MLOps**: NVIDIA Triton, Monarch, TensorRT, ONNX, Volcano, Flyte, LangGraph, DVC
- **Observability**: Grafana, Prometheus, Loki, Alloy, OpenTelemetry, PagerDuty
- **API**: gRPC, gRPC-Gateway, Protobuf, Buf, OpenAPI, Stainless, FastAPI
- **R&D**: PyTorch, Lightning, HuggingFace, Hydra, W&B, NumPy, Pandas, SciPy
- **Data Engineering**: Kafka, RabbitMQ, Celery, Spark, Polars
- **Storage**: PostgreSQL, Redis, MongoDB, S3/R2, Weka, VAST, PGVector
- **Frontend**: React, Next.js, Tailwind