

SUMMARY

Senior Infrastructure/Platform Engineer specializing in distributed systems, GPU infrastructure, and Kubernetes platforms, with 5+ years across AI startups. Currently own Together AI’s RL fine-tuning platform, from training orchestration to multi-cluster GPU scheduling on self-managed clusters. At earlier startups, cut GPU costs by 50%+ with hybrid-cloud infrastructure and shipped production inference at 60k+ req/hour.

EXPERIENCE

-  **Together AI**

Amsterdam, Netherlands

Senior Platform Engineer, Research Platform

July 2025 — Present

 - Designed and now own the RL fine-tuning platform end-to-end, from gRPC/HTTP API and Kafka-driven orchestration to Volcano-based GPU scheduling and Weka/VAST-backed storage, supporting multi-week distributed training runs across up to 256 H100/B200 GPUs (32 nodes)
 - Adopted Meta’s Monarch actor framework into the RL platform, working closely with the PyTorch team on the Kubernetes integration, now powering multi-node GPU training and inference workflows for 7 research engineers
 - Architected the platform’s multi-tier capacity model with best-fit placement across dedicated, reserved, and shared GPU pools, with the latter two autoscaled to demand, eliminating manual capacity planning per customer tier
 - Migrated platform workloads from a click-ops shared cluster to a Terraform-based, Karpenter-autoscaled EKS cluster, eliminating 1-hour manual node scaling and supporting 30+ concurrent per-PR preview environments via ArgoCD ApplicationSet generators
 - Built RL integration and E2E test suites from scratch, running against the live preview environments, and hardened fine-tuning integration tests unlocking dedicated inference for 90+ base models
 - Consolidated the team’s observability from a fragmented setup into a unified Grafana Cloud stack with IaC-managed dashboards and alerting across 9+ K8s clusters, with new clusters now onboardable in minutes
 - Modernized CI/CD, achieving a 4x speedup with BuildKit optimizations and delivering self-hosted GPU-based CI for 90+ engineers via ARC
-  **iClerk**

Remote

MLOps/DevOps Engineer, AI Platform

July 2023 — July 2025 (part-time until March 2024)

 - Built a Terraform-managed platform across 4 environments, combining AWS EKS with Tailscale Operator-connected K3s GPU clusters and ArgoCD-deployed workloads, reducing GPU costs by >50% compared to hyperscaler pricing
 - Established declarative K8s monitoring and alerting in Grafana Cloud across 3 clusters, cutting issue detection from ad hoc to under 1 hour
 - Optimized a speech inference pipeline using WhisperX for ASR, NeMo for diarization, and Demucs for denoising on NVIDIA Triton with per-request model selection, achieving 11x speedup and 5x cost reduction on 1-hour audio inputs
 - Delivered a multi-stage financial document parser using vision LLMs and LangGraph agents on an async Celery pipeline, achieving 70% end-to-end automation and landing a venture fund as the first client
-  **Myna Labs (from the creators of Snap Lenses and Cameos)**

Remote

MLOps Engineer, Speech Synthesis

November 2022 — March 2024

 - Architected and owned a speech inference platform on FastAPI, Celery, and NVIDIA Triton with per-domain priority queues for text-to-speech, voice conversion, and ASR, serving mobile and web apps at 60k+ req/hour
 - Shipped a voice cloning service on Tortoise TTS using LoRA fine-tuning with precomputed speaker latents, cutting per-voice onboarding time from days to under 10 minutes, enabling custom voices in production
 - Replatformed from Docker Compose to Terraform-provisioned GKE with an elastic GPU node pool across staging and production
- Stealth Finance Startup**

Remote

Data Scientist, Quant Finance

September 2021 — July 2022

 - Built a distributed PySpark pipeline that transformed S&P 500 news and analyst reports into daily company-level FinBERT sentiment features
 - Developed ListNet-based learning-to-rank models for long-short equity selection, evaluated with nDCG, Sharpe, and Sortino

EDUCATION

-  **NRU “Higher School of Economics”**

Moscow, Russia

B.S., Computer Science and Finance, summa cum laude

September 2018 — July 2022

 - Ranked in the top 3% of the cohort
 - Passed CFA Level 1 in February 2021, scoring in the top 10% of candidates worldwide

SKILLS

- Languages:** Go, Python, Rust, C++, SQL, TypeScript
- Infrastructure:** Kubernetes, EKS/GKE/K3s, Karpenter, Terraform, Ansible, MAAS, NetBox, AWS/GCP, Docker, Helm, ArgoCD, Nix
- Networking & Security:** Tailscale, Cloudflare, InfiniBand/UFM, OIDC, TLS/mTLS, IAM/RBAC
- ML:** PyTorch, HuggingFace, W&B, NVIDIA Triton, Monarch, vLLM, SGLang, TensorRT, ONNX, Volcano, Flyte, DVC
- Observability:** Grafana, Prometheus, Loki, Alloy, OpenTelemetry, PagerDuty
- API & Data:** gRPC, Protobuf, OpenAPI, Stainless, FastAPI, Kafka, RabbitMQ, Celery
- Storage:** PostgreSQL, Redis, MongoDB, S3, Weka, VAST
- Frontend:** React, Next.js, Tailwind