

SUMMARY

Senior AI Infrastructure / Platform Engineer specializing in distributed systems and GPU infrastructure for production AI. Owns Together AI's RL fine-tuning platform end-to-end, supporting multi-week runs on up to 256 H100 and B200 GPUs. At earlier startups, cut GPU costs by 50%+ with hybrid-cloud clusters and ran speech inference serving 60k+ requests/hour

EXPERIENCE

🌸 Together AI

Amsterdam, Netherlands

Senior Platform Engineer, Post-Training

July 2025 — Present

- Led the design and implementation of Together AI's Tinker-style RL fine-tuning platform from scratch, spanning gRPC/HTTP APIs, asynchronous orchestration, GPU scheduling, and distributed storage, powering multi-week customer training runs on up to 256 H100 and B200 GPUs and the daily workflows of 15+ in-house research engineers
- Drove the adoption of PyTorch Monarch for multi-node GPU training and inference, ran it on Kubernetes before official support, and reported bugs fixed upstream through a publicly announced collaboration with [Meta's PyTorch team](#)
- Designed a multi-tier capacity model with best-fit placement across dedicated, reserved, and shared GPU pools, shipping its first phases: scale-from-zero cluster provisioning and routing of customer workloads to dedicated capacity
- Migrated workloads from a click-ops shared cluster to Terraform-provisioned, Karpenter-autoscaled EKS, eliminating 1-hour manual scaling and enabling 30+ concurrent PR environments with full service stacks deployed via Helm and ArgoCD
- Set the cross-repo testing standard for fine-tuning and RL with a unified test CLI, building RL integration and E2E suites from scratch and hardening fine-tuning suites that validate 90+ base models on live per-PR deployments
- Standardized Grafana Cloud observability across 9+ Kubernetes clusters and deployed ARC-managed GitHub Actions CPU/GPU runners, making CI 4x faster with BuildKit for a 20+ person team and enabling GPU-based testing for 90+ engineers

🏠 iClerk

Remote

MLOps/DevOps Engineer, Multimodal AI

July 2023 — July 2025 (part-time until March 2024)

- Built a Terraform-managed hybrid platform connecting AWS EKS to externally hosted K3s GPU clusters via the Tailscale Operator and deploying workloads through ArgoCD, cutting GPU costs by more than 50% versus hyperscalers
- Established declarative Grafana Cloud monitoring for 3 Kubernetes clusters, cutting issue detection from ad hoc to under 1 hour
- Optimized NVIDIA Triton speech inference with WhisperX/NeMo/Demucs, processing audio 11x faster at 20% of the cost
- Delivered a Celery-based vision-LLM financial parser, automating 70% of reporting and landing a venture fund as its first client

🔥 Myna Labs (from the creators of Snap Lenses, Cameos, and Glam AI)

Remote

MLOps Engineer, Speech AI

November 2022 — March 2024

- Owned a speech platform serving iOS, Android, and web apps at 60k+ req/hour, with priority queues for TTS, VC, and ASR
- Shipped Tortoise TTS voice cloning with LoRA and speaker latents, cutting onboarding from days to under 10 minutes
- Replatformed from Docker Compose to Terraform-provisioned GKE with elastic GPU node pool across staging and production

Stealth Finance Startup (same founders as iClerk)

Remote

Data Scientist, Quantitative Finance

September 2021 — July 2022

- Engineered daily FinBERT sentiment features in PySpark and trained ListNet models for long-short equity ranking

EDUCATION

🎓 Higher School of Economics

Moscow, Russia

B.S., Computer Science and Finance, summa cum laude

September 2018 — July 2022

- Ranked in the top 3% of the cohort
- Passed CFA Level I in the top 10% globally

SKILLS

- **Languages:** Go, Python, Rust, C++, SQL, TypeScript
- **Infrastructure:** K8s, EKS/GKE/K3s, Helm, ArgoCD, Karpenter, Volcano, AWS/GCP, Terraform, Ansible, Nix, MAAS, NetBox
- **APIs & Distributed Systems:** gRPC, Protobuf, OpenAPI, Stainless, FastAPI, Kafka, RabbitMQ, Celery
- **ML Systems:** PyTorch, Monarch, vLLM, SGLang, NVIDIA Triton, TensorRT, Flyte
- **Observability:** Grafana, Prometheus, Loki, Alloy, OpenTelemetry
- **Networking:** InfiniBand, UFM, Tailscale, Cloudflare
- **Storage:** PostgreSQL, MongoDB, Redis, S3, WEKA, VAST