

SUMMARY

Senior Infrastructure/Platform Engineer specializing in distributed systems, GPU infrastructure, and Kubernetes platforms, with 5+ years across AI startups. Currently own Together AI's RL fine-tuning platform, from training orchestration to multi-cluster GPU scheduling on self-managed clusters. At earlier startups, cut GPU costs by 50%+ with hybrid-cloud infrastructure and shipped production inference at 60k+ req/hour.

EXPERIENCE

 **Together AI**

Amsterdam, Netherlands

Senior Platform Engineer, Research Platform

July 2025 — Present

- Designed and now own the RL fine-tuning platform end-to-end, from gRPC/HTTP API and Kafka-driven orchestration to Volcano-based GPU scheduling and Weka/VAST-backed storage, supporting multi-week distributed training runs across up to 256 H100/B200 GPUs (32 nodes)
- Adopted Meta's Monarch actor framework into the RL platform, working closely with the PyTorch team on the Kubernetes integration, now powering multi-node GPU training and inference workflows for 7 research engineers
- Architected the platform's multi-tier capacity model with best-fit placement across dedicated, reserved, and shared GPU pools, with the latter two autoscaled to demand, eliminating manual capacity planning per customer tier
- Migrated platform workloads from a click-ops shared cluster to a Terraform-based, Karpenter-autoscaled EKS cluster, eliminating 1-hour manual node scaling and supporting 30+ concurrent per-PR preview environments via ArgoCD ApplicationSet generators
- Built RL integration and E2E test suites from scratch, running against the live preview environments, and hardened fine-tuning integration tests unlocking dedicated inference for 90+ base models
- Consolidated the team's observability from a fragmented setup into a unified Grafana Cloud stack with IaC-managed dashboards and alerting across 9+ K8s clusters, with new clusters now onboardable in minutes
- Modernized CI/CD, achieving a 4x speedup with BuildKit optimizations and delivering self-hosted GPU-based CI for 90+ engineers via ARC

 **iClerk**

Remote

MLOps/DevOps Engineer, AI Platform

July 2023 — July 2025 (part-time until March 2024)

- Built a Terraform-managed platform across 4 environments, combining AWS EKS with Tailscale Operator-connected K3s GPU clusters and ArgoCD-deployed workloads, reducing GPU costs by >50% compared to hyperscaler pricing
- Established declarative K8s monitoring and alerting in Grafana Cloud across 3 clusters, cutting issue detection from ad hoc to under 1 hour
- Optimized a speech inference pipeline using WhisperX for ASR, NeMo for diarization, and Demucs for denoising on NVIDIA Triton with per-request model selection, achieving 11x speedup and 5x cost reduction on 1-hour audio inputs
- Delivered a multi-stage financial document parser using vision LLMs and LangGraph agents on an async Celery pipeline, achieving 70% end-to-end automation and landing a venture fund as the first client

 **Myna Labs** (from the creators of **Snap Lenses** and **Snap Cameos**)

Remote

MLOps Engineer, Speech Synthesis

November 2022 — March 2024

- Architected and owned a speech inference platform on FastAPI, Celery, and NVIDIA Triton with per-domain priority queues for text-to-speech, voice conversion, and ASR, serving mobile and web apps at 60k+ req/hour
- Shipped a voice cloning service on Tortoise TTS using LoRA fine-tuning with precomputed speaker latents, cutting per-voice onboarding time from days to under 10 minutes, enabling custom voices in production
- Replatformed from Docker Compose to Terraform-provisioned GKE with an elastic GPU node pool across staging and production

Stealth Finance Startup (same founding team as iClerk)

Remote

Data Scientist, Quant Finance

September 2021 — July 2022

- Built a distributed PySpark pipeline that transformed S&P 500 news and analyst reports into daily company-level FinBERT sentiment features
- Developed ListNet-based learning-to-rank models for long-short equity selection, evaluated with nDCG, Sharpe, and Sortino

EDUCATION

 **Higher School of Economics**

Moscow, Russia

B.S., Computer Science and Finance, summa cum laude

September 2018 — July 2022

- Ranked in the top 3% of the cohort
- Passed **CFA Level 1** in February 2021, scoring in the top 10% of candidates worldwide

SKILLS

- Languages:** Go, Python, Rust, C++, SQL, TypeScript
- Infrastructure:** Kubernetes, EKS/GKE/K3s, Helm, ArgoCD, Karpenter, Docker, Terraform, Ansible, AWS/GCP, MAAS, NetBox, Nix
- ML:** PyTorch, Monarch, vLLM, SGLang, NVIDIA Triton, TensorRT, ONNX, Volcano, Flyte, HuggingFace, W&B, DVC
- Networking & Security:** InfiniBand/UFM, Tailscale, Cloudflare, OIDC, TLS/mTLS, IAM/RBAC
- Observability:** Grafana, Prometheus, Loki, Alloy, OpenTelemetry, PagerDuty
- API & Data:** gRPC, Protobuf, OpenAPI, Stainless, FastAPI, Kafka, RabbitMQ, Celery
- Storage:** PostgreSQL, Redis, MongoDB, S3, Weka, VAST