

SUMMARY

Senior Platform Engineer specializing in distributed systems and GPU infrastructure for production AI. Currently leading the architecture of Together AI's RL fine-tuning platform, supporting multi-week runs on up to 256 H100/B200 GPUs. Previously cut GPU costs by 50%+ and built speech inference systems serving 60k+ requests/hour

EXPERIENCE

Together AI

Amsterdam, Netherlands

Senior Platform Engineer, Post-Training

July 2025 — Present

- Lead the architecture of Together AI's RL fine-tuning platform, built from scratch across gRPC/HTTP APIs, asynchronous orchestration, GPU scheduling, and distributed storage, supporting multi-week runs on up to 256 H100/B200 GPUs
- Owned and implemented the end-to-end integration of PyTorch Monarch into the Kubernetes-based RL service, working directly with its developers at Meta to enable multi-node GPU training and inference
- Architected a multi-tier capacity model with best-fit placement across dedicated, reserved, and shared GPU pools, autoscaling reserved and shared capacity to demand, eliminating manual capacity planning by customer tier
- Migrated workloads from a click-ops shared cluster to Terraform-provisioned, Karpenter-autoscaled EKS, eliminating 1-hour manual scaling and enabling 30+ concurrent PR environments with full service stacks deployed via Helm and ArgoCD
- Standardized platform operations across 9+ K8s clusters with codified Grafana Cloud observability and self-hosted ARC runners, speeding CI by 4x and introducing GPU testing for 90+ engineers

iClerk

Remote

MLOps/DevOps Engineer, Multimodal AI

July 2023 — July 2025 (part-time until March 2024)

- Built a Terraform-managed hybrid platform connecting AWS EKS to externally hosted K3s GPU clusters via the Tailscale Operator and deploying workloads through ArgoCD, cutting GPU costs by >50% versus hyperscalers
- Established declarative Grafana Cloud monitoring for 3 K8s clusters, cutting issue detection from ad hoc to under 1 hour
- Optimized NVIDIA Triton speech inference with WhisperX/NeMo/Demucs, processing audio 11x faster at 20% of the cost
- Delivered a Celery-based vision-LLM financial parser, automating 70% of reporting and landing a venture fund as its first client

Myna Labs (from the creators of Snap Lenses and Snap Cameos)

Remote

MLOps Engineer, Speech AI

November 2022 — March 2024

- Owned a speech platform serving iOS, Android, and web apps at 60k+ req/hour, with priority queues for TTS, VC, and ASR
- Shipped Tortoise TTS voice cloning with LoRA and speaker latents, cutting onboarding from days to under 10 minutes
- Replatformed from Docker Compose to Terraform-provisioned GKE with elastic GPU node pool across staging and production

Stealth Finance Startup (same founding team as iClerk)

Remote

Data Scientist, Quantitative Finance

September 2021 — July 2022

- Engineered daily FinBERT sentiment features in PySpark and trained ListNet models for long-short equity ranking

EDUCATION

Higher School of Economics

Moscow, Russia

B.S., Computer Science and Finance, summa cum laude

September 2018 — July 2022

- Ranked in the top 3% of the cohort
- Passed CFA Level I in the top 10% globally

SKILLS

- **Languages:** Go, Python, Rust, C++, SQL, TypeScript
- **Infrastructure:** K8s, EKS/GKE/K3s, Helm, ArgoCD, Karpenter, AWS/GCP, Terraform, Ansible, MAAS, NetBox, Nix
- **ML:** PyTorch, Monarch, vLLM, SGLang, NVIDIA Triton, TensorRT, Volcano, Flyte, W&B
- **Networking & Security:** InfiniBand/UFM, Tailscale, Cloudflare, Dex, Okta
- **Observability:** Grafana, Prometheus, Loki, Alloy, OpenTelemetry, PagerDuty
- **API & Data:** gRPC, Protobuf, OpenAPI, Stainless, FastAPI, Kafka, RabbitMQ, Celery
- **Storage:** PostgreSQL, Redis, MongoDB, S3, Weka, VAST