

EXPERIENCE

🌸 Together AI

Senior Platform Engineer, Post-Training

Amsterdam, Netherlands

July 2025 — Present

- Led the architecture of the Tinker-compatible [custom training platform](#) and built the service end to end. It reached customer beta in 10 weeks and now runs multi-week RL experiments on up to 144 H100/B200 GPUs.
- Unified managed fine-tuning and custom training on a multi-cluster pool of 800+ H100/B200 GPUs, routing tenants to dedicated or shared capacity so idle GPUs can be sold as [preemptible compute](#).
- Took pre-release PyTorch Monarch into production as the training and inference runtime that 20+ research engineers now develop on. Ran it on Kubernetes before official support, working directly with [Meta's PyTorch team](#).
- Migrated all post-training services from a shared, hand-managed cluster to Terraform-provisioned EKS with Karpenter autoscaling, a first at Together that became the template for other teams. Capacity scales with demand, not by hand, so 40+ per-PR preview environments run integration and end-to-end tests at no idle cost.
- Moved CI onto self-hosted GitHub Actions runners on Kubernetes with ARC. BuildKit caching made the team's CI 4x faster, and GPU runners opened GPU testing to 90+ engineers across the research organization.

💻 iClerk

MLOps/DevOps Engineer, Multimodal AI

Remote

July 2023 — July 2025 (part-time to March 2024)

- Designed and ran the hybrid platform: EKS for services, a K3s GPU cluster on low-cost providers, and a Tailscale mesh between them. GPU costs came in more than 50% below AWS.
- Replaced vanilla Whisper transcription with an NVIDIA Triton pipeline: WhisperX, NeMo diarization, and Demucs denoising behind one BLS orchestrator. The pipeline ran 11x faster at 20% of the cost.
- Built a vision-LLM pipeline for parsing 10-K/10-Q filings using LangGraph and Celery, automating ~70% of the analyst workflow and helping land a venture fund as one of the company's first customers.
- Owned observability across 3 Kubernetes clusters: Grafana Cloud with Alloy collectors, everything from the stack itself to dashboards and alerts defined in Terraform. Wrote up the approach as a [public guide](#).

🦋 Myna Labs (from the creators of Snap Lenses/Cameos and Glam AI)

MLOps Engineer, Speech AI

Remote

November 2022 — March 2024

- Owned the speech backend serving iOS, Android, and web apps at 60k+ requests/hour, with separate recognition, synthesis, and voice-conversion queues that prioritized requests by the customer's subscription tier.
- Shipped Tortoise TTS custom voices to the consumer apps. Instead of retraining per voice, an NVIDIA Triton BLS pipeline extracted a LoRA adapter and speaker latents, so a new voice was ready in under 10 minutes, not days.

💎 AgBlox (same founders as iClerk)

Data Scientist, Quantitative Finance

Remote

September 2021 — July 2022

- Built FinBERT sentiment features and ListNet equity rankers in PySpark for a long-short strategy.

EDUCATION

🎓 Higher School of Economics

B.S., Computer Science and Finance, summa cum laude

Moscow, Russia

September 2018 — July 2022

- Ranked in the top 3% of the cohort.
- Passed CFA Level I in the top 10% globally.

SKILLS

- **Languages:** Go, Python, Rust, C++, SQL, TypeScript
- **Infrastructure:** Kubernetes (EKS, GKE, K3s), ArgoCD, Karpenter, Volcano, AWS, GCP, Terraform, Ansible, Nix
- **Distributed Systems & APIs:** gRPC, Protobuf, Kafka, RabbitMQ, OpenAPI, Stainless
- **ML Systems:** PyTorch, Monarch, vLLM, SGLang, NVIDIA Triton, TensorRT, Flyte
- **Observability:** Grafana Cloud, Prometheus, Loki, OpenTelemetry, Alloy
- **Networking & Security:** InfiniBand, Tailscale, Cloudflare, OIDC (Okta, Dex)
- **Data & Storage:** PostgreSQL, Redis, S3, MongoDB